

# Part 3: Machine Learning – Predictive Modeling Objective

Build predictive models to forecast demand and identify high-tipping scenarios.

## Tasks & Models

### 1. Weekly Trip Demand Forecasting

Goal: Predict number of trips per week using historical trip counts and seasonal patterns.

### 2. High-Tip Prediction Model

Goal: Predict if a trip will have a tip\_amount above the average using trip\_distance, boroughs, and time-based features.

### 3. Route-Based Revenue Forecasting (Optional)

Goal: Forecast total\_amount for a specific pickup\_borough–dropoff\_borough pair in future months.

```
# Basic imports
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score,
classification_report
from sklearn.ensemble import RandomForestRegressor,
RandomForestClassifier

import pandas as pd
import glob

# folder
folder_path = r"C:/Users/parth/OneDrive/Desktop/Trips/"

# Load CSV
all_files = glob.glob(folder_path + "*.csv")

# Show files
print("Files found:", all_files)

# Read and combine
df_list = []
for file in all_files:
    print("Loading:", file)
```

```

temp_df = pd.read_csv(file)
df_list.append(temp_df)

# Combine all years
df = pd.concat(df_list, ignore_index=True)

# Show preview
df.head(), df.shape

Files found: ['C:/Users/parth/OneDrive/Desktop/Trips\\
2017_trimmed.csv', 'C:/Users/parth/OneDrive/Desktop/Trips\\
2018_trimmed.csv', 'C:/Users/parth/OneDrive/Desktop/Trips\\
2019_trimmed.csv', 'C:/Users/parth/OneDrive/Desktop/Trips\\
2020_trimmed.csv']
Loading: C:/Users/parth/OneDrive/Desktop/Trips\2017_trimmed.csv
Loading: C:/Users/parth/OneDrive/Desktop/Trips\2018_trimmed.csv
Loading: C:/Users/parth/OneDrive/Desktop/Trips\2019_trimmed.csv
Loading: C:/Users/parth/OneDrive/Desktop/Trips\2020_trimmed.csv

(   VendorID      lpep_pickup_datetime      lpep_dropoff_datetime  \
0        2.0  2017-01-04 18:03:23.000  2017-01-04 18:10:41.000
1        2.0  2017-02-21 14:36:40.000  2017-02-21 14:44:06.000
2        2.0  2017-03-09 08:53:53.000  2017-03-09 08:59:02.000
3        2.0  2017-12-05 20:15:50.000  2017-12-05 20:18:26.000
4        2.0  2017-07-12 14:45:33.000  2017-07-12 14:50:52.000

      store_and_fwd_flag  RatecodeID  PULocationID  DOLocationID
passenger_count \
0                  N       1.0           33            52
1.0
1                  N       1.0           25            97
1.0
2                  N       1.0           41           166
1.0
3                  N       1.0          260           260
5.0
4                  N       1.0           17            17
1.0

      trip_distance  fare_amount  extra  mta_tax  tip_amount
tolls_amount \
0             0.96        6.5    1.0     0.5      1.66
0.0
1             1.12        6.5    0.0     0.5      2.19
0.0
2             0.95        6.0    0.0     0.5      1.36
0.0
3             0.55        4.0    0.5     0.5      1.00
0.0
4             0.63        5.5    0.0     0.5      0.00

```

```
0.0
```

```
    improvement_surcharge  total_amount  payment_type  trip_type \
0                  0.3        9.96          1.0         1.0
1                  0.3        9.49          1.0         1.0
2                  0.3        8.16          1.0         1.0
3                  0.3        6.30          1.0         1.0
4                  0.3        6.30          2.0         1.0

    congestion_surcharge
0                   NaN
1                   NaN
2                   NaN
3                   NaN
4                   NaN  ,
(4000000, 19))

# Ensure datetime column is in datetime format
df["lpep_pickup_datetime"] =
pd.to_datetime(df["lpep_pickup_datetime"])

# Create a week start date column
df["week_start"] =
df["lpep_pickup_datetime"].dt.to_period("W").apply(lambda r:
r.start_time)

weekly = (
    df.groupby("week_start")
    .size()
    .reset_index(name="trip_count")
    .sort_values("week_start")
)
weekly.head(), weekly.tail()

(  week_start  trip_count
0 2008-12-29      8
1 2010-09-20      3
2 2016-12-26    318
3 2017-01-02   1956
4 2017-01-09   2170,
    week_start  trip_count
207 2020-11-30    1193
208 2020-12-07    1249
209 2020-12-14    1053
210 2020-12-21     980
211 2020-12-28     642)

# Create lag features: previous 1, 2, 3 weeks' trip counts
for lag in [1, 2, 3]:
```

```

weekly[f"lag_{lag}"] = weekly["trip_count"].shift(lag)

# Drop first few rows with NaNs from lagging
weekly_ml = weekly.dropna().reset_index(drop=True)

X = weekly_ml[["lag_1", "lag_2", "lag_3"]]
y = weekly_ml["trip_count"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False
)

model_weekly = RandomForestRegressor(random_state=42)
model_weekly.fit(X_train, y_train)

y_pred = model_weekly.predict(X_test)

print("Weekly demand forecasting RMSE:", mean_squared_error(y_test,
y_pred, squared=False))
print("R2:", r2_score(y_test, y_pred))

# Plot actual vs predicted
plt.figure()
plt.plot(y_test.values, label="Actual")
plt.plot(y_pred, label="Predicted")
plt.title("Weekly Trip Demand – Actual vs Predicted")
plt.xlabel("Test Weeks")
plt.ylabel("Trips")
plt.legend()
plt.show()

-----
-----
KeyError                                     Traceback (most recent call
last)
File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\
pandas\core\indexes\base.py:3812, in Index.get_loc(self, key)
    3811 try:
-> 3812     return self._engine.get_loc(casted_key)
    3813 except KeyError as err:

File pandas/_libs/index.pyx:167, in
pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/index.pyx:175, in
pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/index_class_helper.pxi:245, in
pandas._libs.index.MaskedUInt32Engine._check_type()

KeyError: 'trip_count'
```

The above exception was the direct cause of the following exception:

```
KeyError Traceback (most recent call
last)
Cell In[56], line 3
    1 # Create lag features: previous 1, 2, 3 weeks' trip counts
    2 for lag in [1, 2, 3]:
----> 3     weekly[f"lag_{lag}"] = weekly["trip_count"].shift(lag)
    5 # Drop first few rows with NaNs from lagging
    6 weekly_ml = weekly.dropna().reset_index(drop=True)

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\
pandas\core\series.py:1130, in Series.__getitem__(self, key)
    1127     return self._values[key]
    1129 elif key_is_scalar:
-> 1130     return self._get_value(key)
    1132 # Convert generator to list before going through hashable part
    1133 # (We will iterate through the generator there to check for
slices)
    1134 if is_iterator(key):

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\
pandas\core\series.py:1246, in Series._get_value(self, label,
takeable)
    1243     return self._values[label]
    1245 # Similar to Index.get_value, but we do not fall back to
positional
-> 1246 loc = self.index.get_loc(label)
    1248 if is_integer(loc):
    1249     return self._values[loc]

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\
pandas\core\indexes\base.py:3819, in Index.get_loc(self, key)
    3814     if isinstance(casted_key, slice) or (
    3815         isinstance(casted_key, abc.Iterable)
    3816         and any(isinstance(x, slice) for x in casted_key)
    3817     ):
    3818         raise InvalidIndexError(key)
-> 3819     raise KeyError(key) from err
    3820 except TypeError:
    3821     # If we have a listlike key, _check_indexing_error will
raise
    3822     # InvalidIndexError. Otherwise we fall through and re-
raise
    3823     # the TypeError.
    3824     self._check_indexing_error(key)

KeyError: 'trip_count'
```

```

# =====
# □ Weekly Aggregation (using previous approach)
# =====
df['lpep_pickup_datetime'] =
pd.to_datetime(df['lpep_pickup_datetime'])

df['week_start'] =
df['lpep_pickup_datetime'].dt.to_period("W").apply(lambda r:
r.start_time)

weekly = (
    df.groupby('week_start')
    .size()
    .reset_index(name='trip_count')
    .sort_values('week_start')
)
print("Weekly data:")
weekly.head()

Weekly data:

   week_start  trip_count
0  2008-12-29          8
1  2010-09-20          3
2  2016-12-26        318
3  2017-01-02       1956
4  2017-01-09       2170

# xyz
# Create lag features
# xyz
for lag in [1, 2, 3]:
    weekly[f'lag_{lag}'] = weekly['trip_count'].shift(lag)

weekly_ml = weekly.dropna().reset_index(drop=True)

X = weekly_ml[['lag_1', 'lag_2', 'lag_3']]
y = weekly_ml['trip_count']

print(weekly_ml.head())

   week_start  trip_count  lag_1  lag_2  lag_3
0  2017-01-02      1956.0  318.0    3.0    8.0
1  2017-01-09      2170.0  1956.0  318.0    3.0
2  2017-01-16      2006.0  2170.0  1956.0  318.0
3  2017-01-23      2211.0  2006.0  2170.0  1956.0
4  2017-01-30      2262.0  2211.0  2006.0  2170.0

# xyz
#Train RandomForest model (previous code style)

```

```

# xyz
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False
)

model_weekly = RandomForestRegressor(random_state=42)
model_weekly.fit(X_train, y_train)

y_pred = model_weekly.predict(X_test)

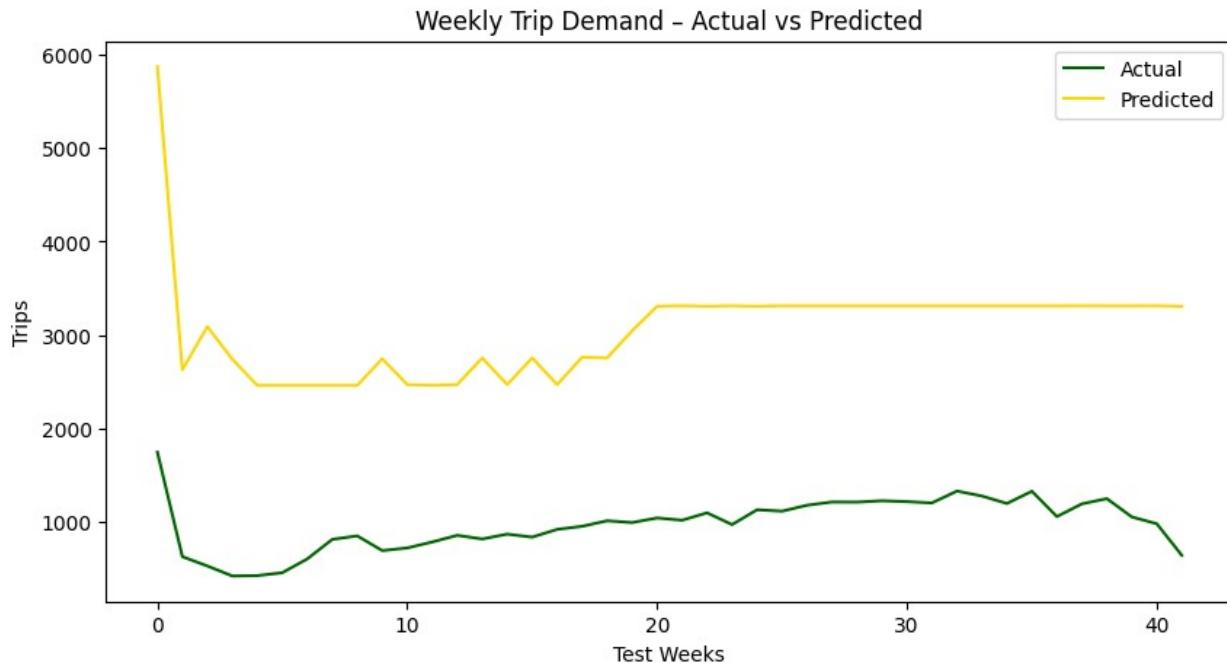
rmse = mean_squared_error(y_test, y_pred) ** 0.5 # works for all
sklearn versions
print("Weekly demand forecasting RMSE:", rmse)
print("R²:", r2_score(y_test, y_pred))

Weekly demand forecasting RMSE: 2129.342294309051
R²: -57.46218052500757

# =====
# □ Plot Actual vs Predicted
# =====
plt.figure()
plt.figure(figsize=(10,5))
plt.plot(y_test.values, color="#006400", label="Actual") # Dark
Green
plt.plot(y_pred, color="#FFD700", label="Predicted") # Dark
Yellow (Gold)
plt.title("Weekly Trip Demand – Actual vs Predicted")
plt.xlabel("Test Weeks")
plt.ylabel("Trips")
plt.legend()
plt.show()

<Figure size 640x480 with 0 Axes>

```



```

# Add year column if not already added
df['year'] = df['lpep_pickup_datetime'].dt.year

unique_years = sorted(df['year'].unique())

for yr in unique_years:
    df_year = df[df['year'] == yr]

    # Recreate weekly grouping for each year
    weekly =
df_year.groupby(df_year['lpep_pickup_datetime'].dt.isocalendar().week)
.size()

    # Split into train-test for each year
    X = weekly.index.values.reshape(-1, 1)
    y = weekly.values

    from sklearn.model_selection import train_test_split
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

    # Fit model
    from sklearn.ensemble import RandomForestRegressor
    model = RandomForestRegressor()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

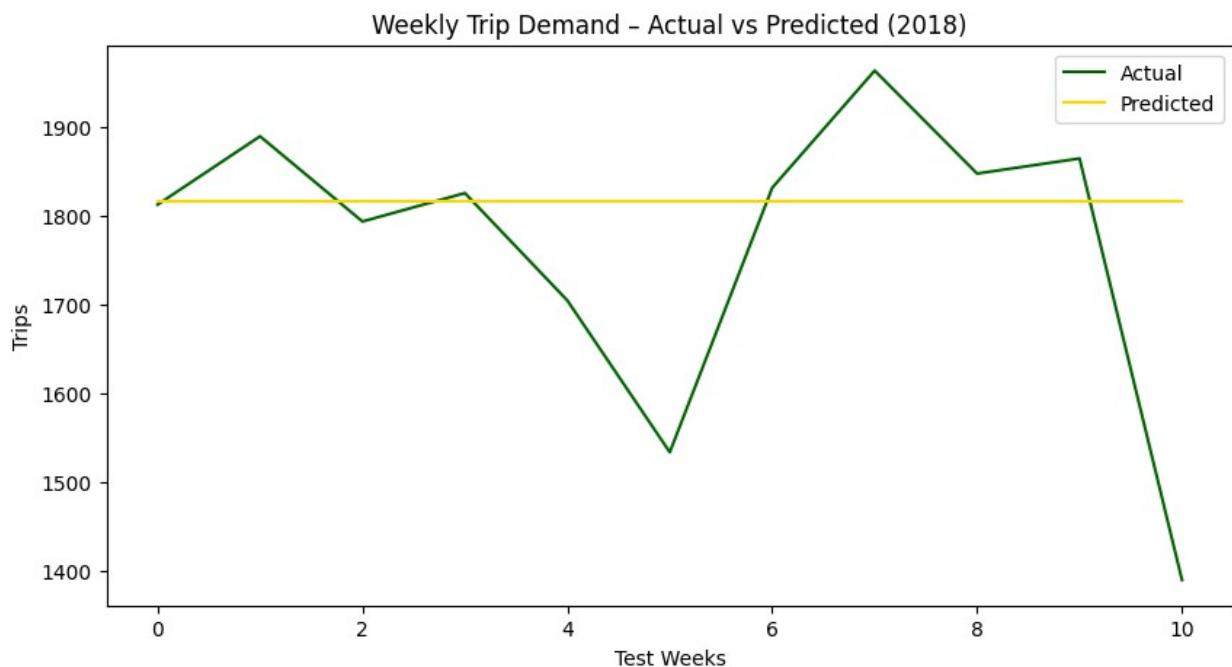
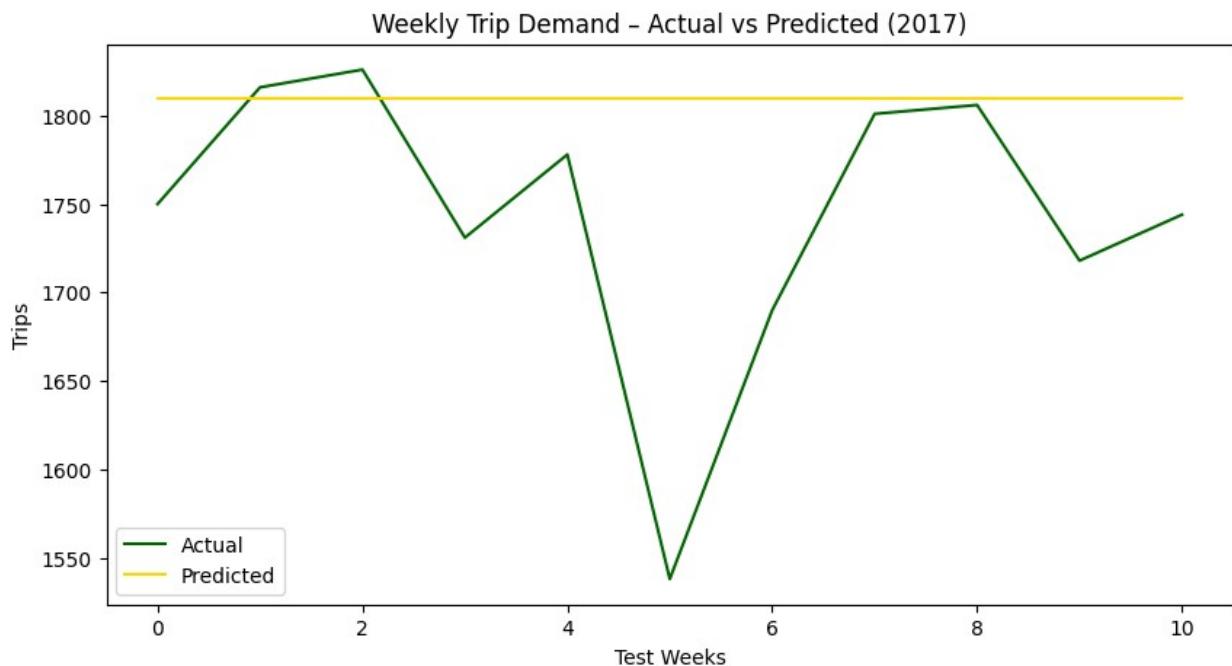
    # Plot
    plt.figure(figsize=(10,5))

```

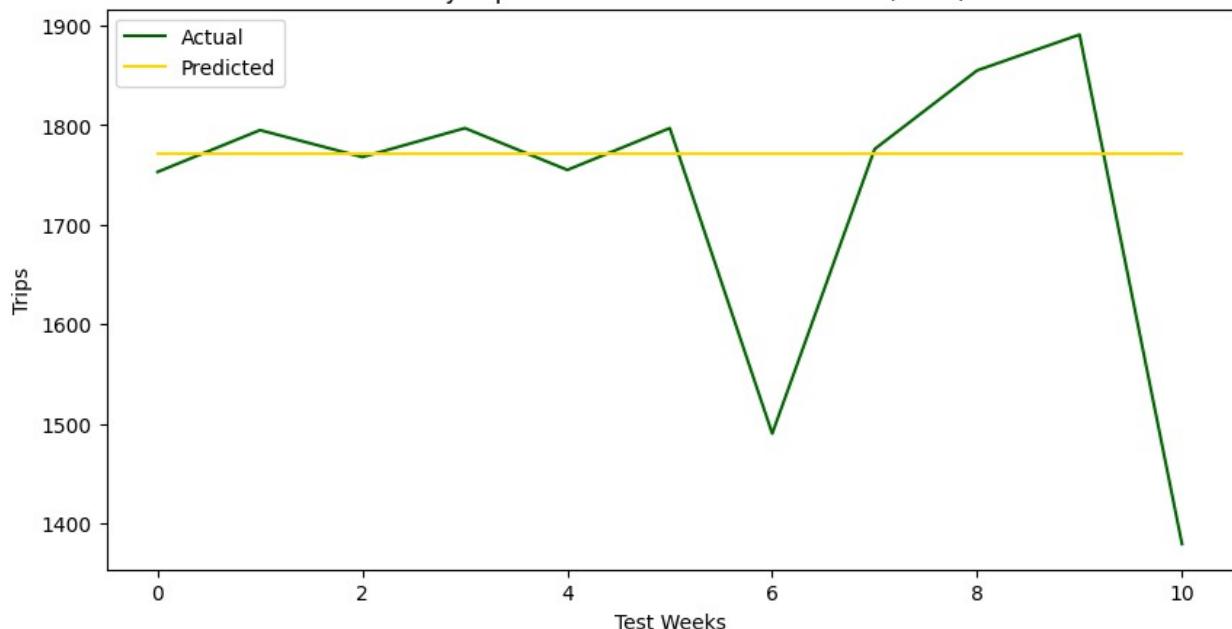
```

plt.plot(y_test, color="#006400", label="Actual")      # Dark green
plt.plot(y_pred, color="#FFD700", label="Predicted") # Gold
plt.title(f"Weekly Trip Demand – Actual vs Predicted ({yr})")
plt.xlabel("Test Weeks")
plt.ylabel("Trips")
plt.legend()
plt.show()

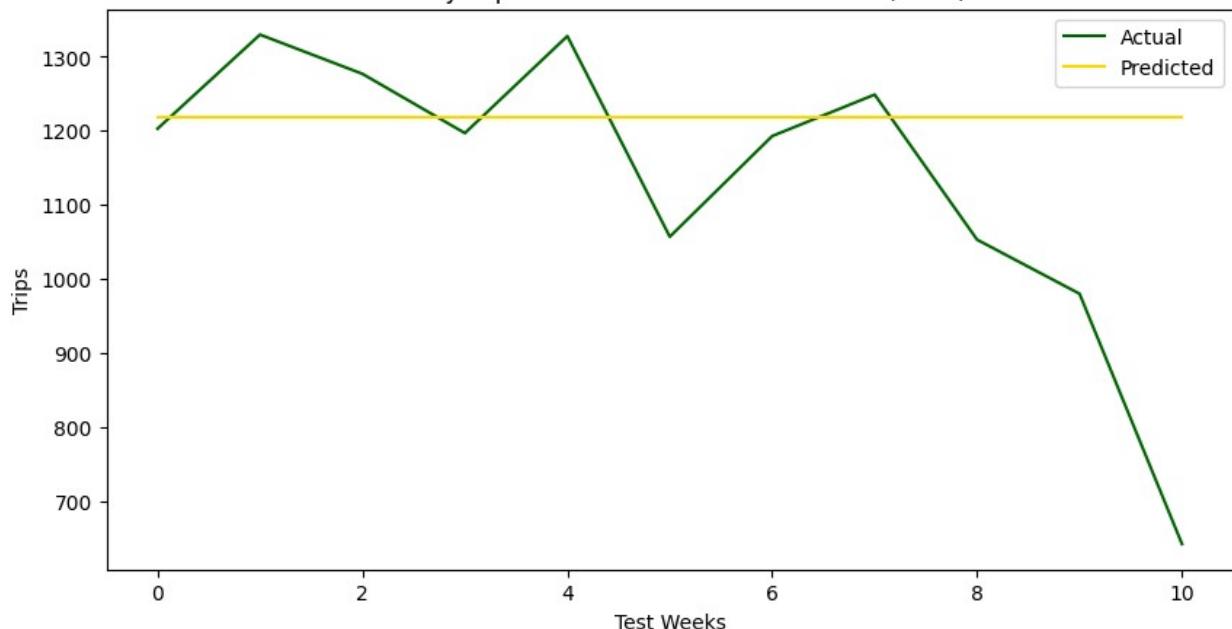
```



Weekly Trip Demand - Actual vs Predicted (2019)



Weekly Trip Demand - Actual vs Predicted (2020)



```
# =====
# Create High-Tip Target
# =====
df2 = df.copy()
df2 = df2[df2['tip_amount'] >= 0] # remove invalid values

avg_tip = df2['tip_amount'].mean()
print("Average tip amount:", avg_tip)
```

```

df2['high_tip'] = (df2['tip_amount'] > avg_tip).astype(int)

Average tip amount: 1.1060786931477258

# =====
# Create Model Features
# xyz
df2['hour'] = df2['lpep_pickup_datetime'].dt.hour
df2['day_of_week'] = df2['lpep_pickup_datetime'].dt.dayofweek # Monday = 0
df2['month'] = df2['lpep_pickup_datetime'].dt.month

# Columns that must exist:
# trip_distance, pickup_borough, dropoff_borough

df2.columns.tolist()

['VendorID',
 'lpep_pickup_datetime',
 'lpep_dropoff_datetime',
 'store_and_fwd_flag',
 'RatecodeID',
 'PULocationID',
 'DOLocationID',
 'passenger_count',
 'trip_distance',
 'fare_amount',
 'extra',
 'mta_tax',
 'tip_amount',
 'tolls_amount',
 'improvement_surcharge',
 'total_amount',
 'payment_type',
 'trip_type',
 'congestion_surcharge',
 'week_start',
 'high_tip',
 'hour',
 'day_of_week',
 'month']

# =====
# Select Columns for Model (using existing columns)
# xyz
cat_cols = ['PULocationID', 'DOLocationID']      # categorical
features
num_cols = ['trip_distance', 'passenger_count',
            'hour', 'day_of_week', 'month']      # numeric / time
features

```

```

# Build modeling dataframe
df_model = pd.get_dummies(
    df2[cat_cols + num_cols + ['high_tip']],
    columns=cat_cols,
    drop_first=True
)

X = df_model.drop("high_tip", axis=1)
y = df_model["high_tip"]

print("Model feature shape:", X.shape)
X.head()

Model feature shape: (399984, 520)

      trip_distance  passenger_count  hour  day_of_week  month
PUlocationID_3 \
0            0.96                 1.0     18             2      1
False
1            1.12                 1.0     14             1      2
False
2            0.95                 1.0      8             3      3
False
3            0.55                 5.0    20             1     12
False
4            0.63                 1.0     14             2      7
False

      PUlocationID_4  PUlocationID_5  PUlocationID_6  PUlocationID_7  ...
\
0            False        False        False        False        ...
1            False        False        False        False        ...
2            False        False        False        False        ...
3            False        False        False        False        ...
4            False        False        False        False        ...

      DOLocationID_256  DOLocationID_257  DOLocationID_258
DOLocationID_259 \
0            False        False        False
False
1            False        False        False
False
2            False        False        False
False
3            False        False        False

```

```

False
4      False      False      False
False

    DOLocationID_260  DOLocationID_261  DOLocationID_262
DOLocationID_263 \
0      False      False      False
False
1      False      False      False
False
2      False      False      False
False
3      True       False      False
False
4      False      False      False
False

    DOLocationID_264  DOLocationID_265
0      False      False
1      False      False
2      False      False
3      False      False
4      False      False

[5 rows x 520 columns]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

clf = RandomForestClassifier(
    n_estimators=150,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))

Accuracy: 0.7094766053727015

Classification Report:

          precision    recall   f1-score   support
          0        0.75     0.83     0.79     51808

```

	1	0.61	0.49	0.54	28189
accuracy				0.71	79997
macro avg		0.68	0.66	0.67	79997
weighted avg		0.70	0.71	0.70	79997

```

# =====
# Train RandomForest Classifier
# xyz
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

clf = RandomForestClassifier(
    n_estimators=150,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

# =====
# Model Evaluation
# xyz
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))

Accuracy: 0.7094766053727015

Classification Report:

      precision    recall  f1-score   support

          0       0.75      0.83      0.79     51808
          1       0.61      0.49      0.54     28189

accuracy                           0.71      79997
macro avg                           0.68      0.66      0.67      79997
weighted avg                          0.70      0.71      0.70      79997

!pip install seaborn

```

```
Requirement already satisfied: seaborn in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from seaborn) (2.3.2)
Requirement already satisfied: pandas>=1.2 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from seaborn) (2.3.1)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from seaborn) (3.10.6)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.3)
Requirement already satisfied: cycler>=0.10 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.59.2)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.9)
Requirement already satisfied: packaging>=20.0 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (25.0)
Requirement already satisfied: pillow>=8 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.3)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: six>=1.5 in c:\users\parth\appdata\local\programs\python\python313\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.17.0)

WARNING: Ignoring invalid distribution ~treamlit (C:\Users\parth\AppData\Local\Programs\Python\Python313\Lib\site-packages)
WARNING: Ignoring invalid distribution ~treamlit (C:\Users\parth\AppData\Local\Programs\Python\Python313\Lib\site-packages)
WARNING: Ignoring invalid distribution ~treamlit (C:\Users\parth\
```

```

AppData\Local\Programs\Python\Python313\Lib\site-packages)

[notice] A new release of pip is available: 25.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip

# Feature Importance
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Use your trained RandomForestClassifier
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1]    # sort descending
features_list = X.columns

feat_imp_df = pd.DataFrame({
    'Feature': features_list[indices],
    'Importance': importances[indices]
})

feat_imp_df.head(10)

      Feature  Importance
0  trip_distance  0.243412
1          hour  0.137212
2        month  0.124673
3  day_of_week  0.096598
4  passenger_count  0.053187
5  PULocationID_66  0.008116
6  DOLocationID_138  0.004482
7  PULocationID_255  0.004215
8  PULocationID_33  0.003505
9  PULocationID_74  0.003426

# Top 10 Most Important Features
# xyz
plt.figure(figsize=(10,6))

sns.barplot(
    data = feat_imp_df.head(10),
    x = 'Importance',
    y = 'Feature',
    palette = 'Blues_r',
    dodge = False
)

plt.title("Top 10 Features Influencing High-Tip Prediction",
          fontsize=14, weight='bold')

plt.xlabel("Feature Importance (Weight)", fontsize=12)

```

```

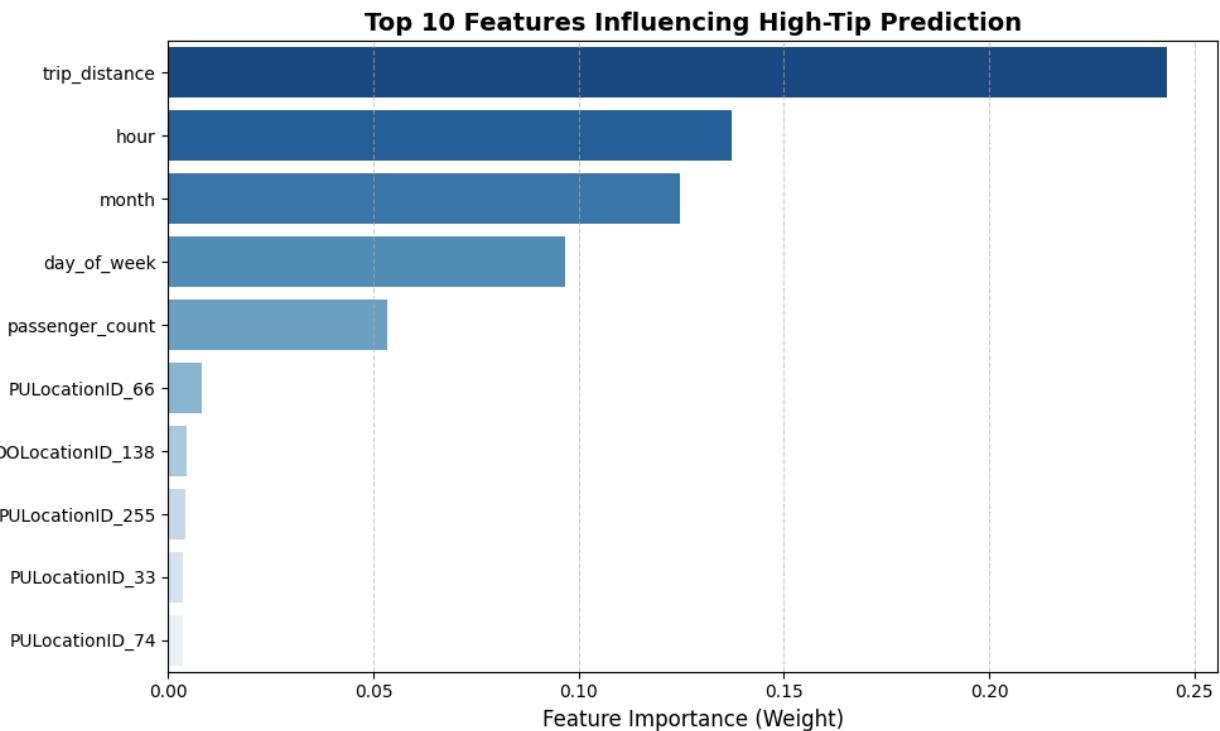
plt.ylabel("")
plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

C:\Users\parth\AppData\Local\Temp\ipykernel_15004\2225236825.py:6:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `y` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(

```



```

import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

# Feature Importance Calculation
# xyz
import numpy as np
import pandas as pd

importances = clf.feature_importances_          # <- IMPORTANT
indices = np.argsort(importances)[::-1]
features_list = X.columns

```

```

feat_imp_df = pd.DataFrame({
    'Feature': features_list[indices],
    'Importance': importances[indices]
})

feat_imp_df.head()

      Feature  Importance
0  trip_distance    0.243412
1          hour     0.137212
2         month     0.124673
3  day_of_week     0.096598
4  passenger_count    0.053187

cat_cols = ['PULocationID', 'DOLocationID']
num_cols = ['trip_distance', 'passenger_count', 'hour', 'day_of_week',
'month']

df_model = pd.get_dummies(
    df2[cat_cols + num_cols + ['high_tip']],
    columns=cat_cols,
    drop_first=True
)

X = df_model.drop("high_tip", axis=1)
y = df_model["high_tip"]

print(X.shape, y.shape)

(399984, 520) (399984,)

# High-Tip Prediction Model + Top 10 Feature Importance
# xyz

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Build feature matrix X and target y
cat_cols = ['PULocationID', 'DOLocationID'] # categorical
num_cols = ['trip_distance', 'passenger_count',
            'hour', 'day_of_week', 'month'] # numeric / time

df_model = pd.get_dummies(
    df2[cat_cols + num_cols + ['high_tip']],

```

```

        columns=cat_cols,
        drop_first=True
    )

X = df_model.drop('high_tip', axis=1)
y = df_model['high_tip']

print("Feature matrix shape:", X.shape)
print("Target distribution:\n", y.value_counts(normalize=True))

❷ 2 Train / Test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

❸ 3 Train RandomForest classifier
clf = RandomForestClassifier(
    n_estimators=150,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

❹ 4 Evaluation
print("\nAccuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))

❺ 5 Feature importance calculation
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1]
features_list = X.columns

feat_imp_df = pd.DataFrame({
    'Feature': features_list[indices],
    'Importance': importances[indices]
})

print("\nTop 10 features by importance:\n", feat_imp_df.head(10))

❻ 6 Plot Top 10 Most Important Features
plt.figure(figsize=(10, 6))
sns.barplot(
    data=feat_imp_df.head(10),
    x='Importance',
    y='Feature',
    dodge=False
)

```

```
)  
  
plt.title("Top 10 Features Influencing High-Tip Prediction",  
          fontsize=14, weight='bold')  
plt.xlabel("Feature Importance (Weight)", fontsize=12)  
plt.ylabel("")  
plt.grid(axis='x', linestyle='--', alpha=0.6)  
plt.tight_layout()  
plt.show()
```

Feature matrix shape: (399984, 520)

Target distribution:

```
high_tip  
0    0.647626  
1    0.352374
```

Name: proportion, dtype: float64

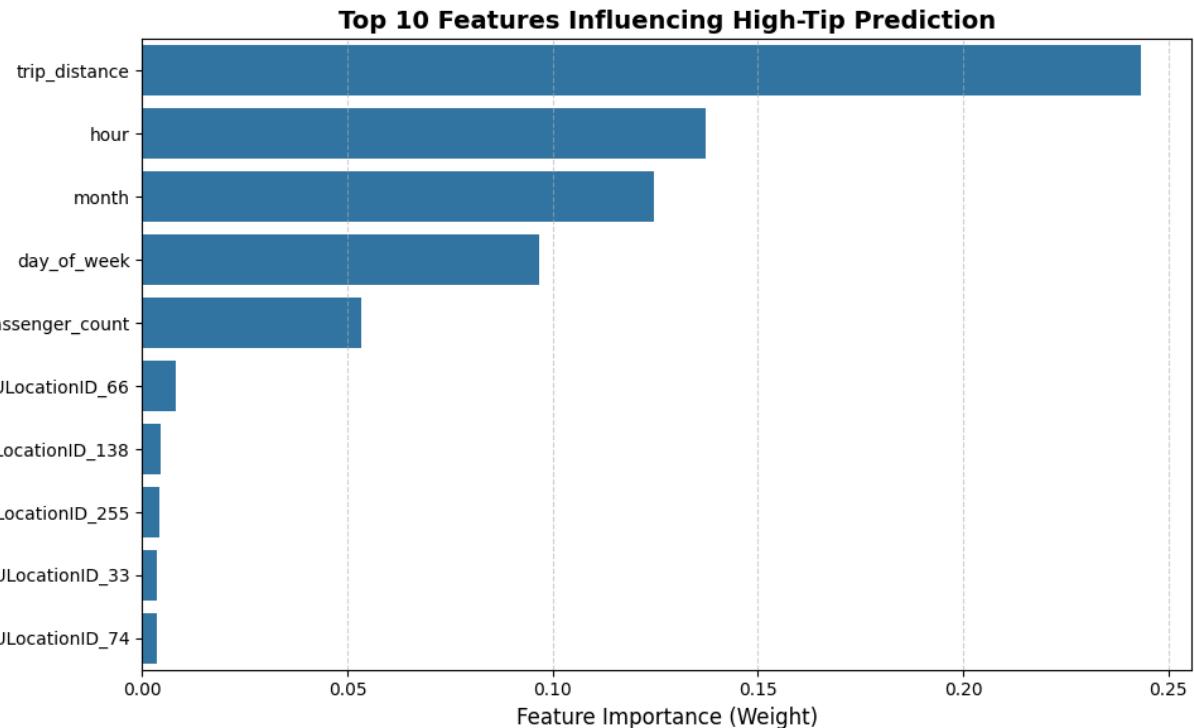
Accuracy: 0.7094641049039339

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.83	0.79	51808
1	0.61	0.49	0.54	28189
accuracy			0.71	79997
macro avg	0.68	0.66	0.67	79997
weighted avg	0.70	0.71	0.70	79997

Top 10 features by importance:

	Feature	Importance
0	trip_distance	0.243412
1	hour	0.137212
2	month	0.124673
3	day_of_week	0.096598
4	passenger_count	0.053187
5	PULocationID_66	0.008116
6	DOLocationID_138	0.004482
7	PULocationID_255	0.004215
8	PULocationID_33	0.003505
9	PULocationID_74	0.003426



### *# Route-Based Monthly Revenue Forecasting*

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# 1 Ensure correct datetime format
df['lpep_pickup_datetime'] =
pd.to_datetime(df['lpep_pickup_datetime'], errors='coerce')

# 2 Identify top route (highest trip count)
route_counts = (
    df.groupby(['PUlocationID', 'DOlocationID'])
    .size()
    .reset_index(name='trip_count')
    .sort_values('trip_count', ascending=False)
)

print("Top 10 busiest routes:")
display(route_counts.head(10))

# Automatically pick the most frequent route

```

```

top_route = route_counts.iloc[0]
pu_id = int(top_route['PULocationID'])
do_id = int(top_route['DOLocationID'])

print(f"\nUsing Route: PULocationID={pu_id} → DOLocationID={do_id}")

❸ 3 Filter data for selected route
route_df = df[(df['PULocationID'] == pu_id) &
               (df['DOLocationID'] == do_id)].copy()

print("Total records for selected route:", len(route_df))

❹ 4 Create monthly revenue
route_df['year_month'] =
    route_df['lpep_pickup_datetime'].dt.to_period('M').dt.to_timestamp()

monthly_rev = (
    route_df.groupby('year_month')['total_amount']
        .sum()
        .reset_index(name='monthly_revenue')
        .sort_values('year_month')
)
print("\nMonthly revenue preview:")
display(monthly_rev.head())

❺ 5 Create lag features
monthly_rev['lag_1'] = monthly_rev['monthly_revenue'].shift(1)
monthly_rev['lag_2'] = monthly_rev['monthly_revenue'].shift(2)

# Drop NaNs
monthly_ml = monthly_rev.dropna().reset_index(drop=True)

X = monthly_ml[['lag_1', 'lag_2']]
y = monthly_ml['monthly_revenue']

❻ 6 Train / Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, shuffle=False # keep time order
)

❼ 7 Train RandomForestRegressor
model = RandomForestRegressor(
    n_estimators=150,
    random_state=42
)

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

```

```
❸ 8 Performance Metrics
rmse = mean_squared_error(y_test, y_pred) ** 0.5
r2 = r2_score(y_test, y_pred)

print("\n▣ Route Revenue Forecasting Results")
print("RMSE:", rmse)
print("R² Score:", r2)

❹ 9 Plot Actual vs Predicted
plt.figure(figsize=(10,5))
plt.plot(y_test.values, label="Actual", color="#006400")          # Dark green
plt.plot(y_pred, label="Predicted", color="#FFD700")            # Dark yellow
plt.title(f"Monthly Revenue – Route {pu_id} → {do_id}")
plt.xlabel("Test Months")
plt.ylabel("Total Amount ($)")
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

## Top 10 busiest routes:

PULocationID	DOLocationID	trip_count
7029	75	4594
130	7	4545
6822	74	3830
3473	41	3789
9340	95	3358
6821	74	3056
8039	82	3012
3672	42	3001
3472	41	2934
3496	41	2809

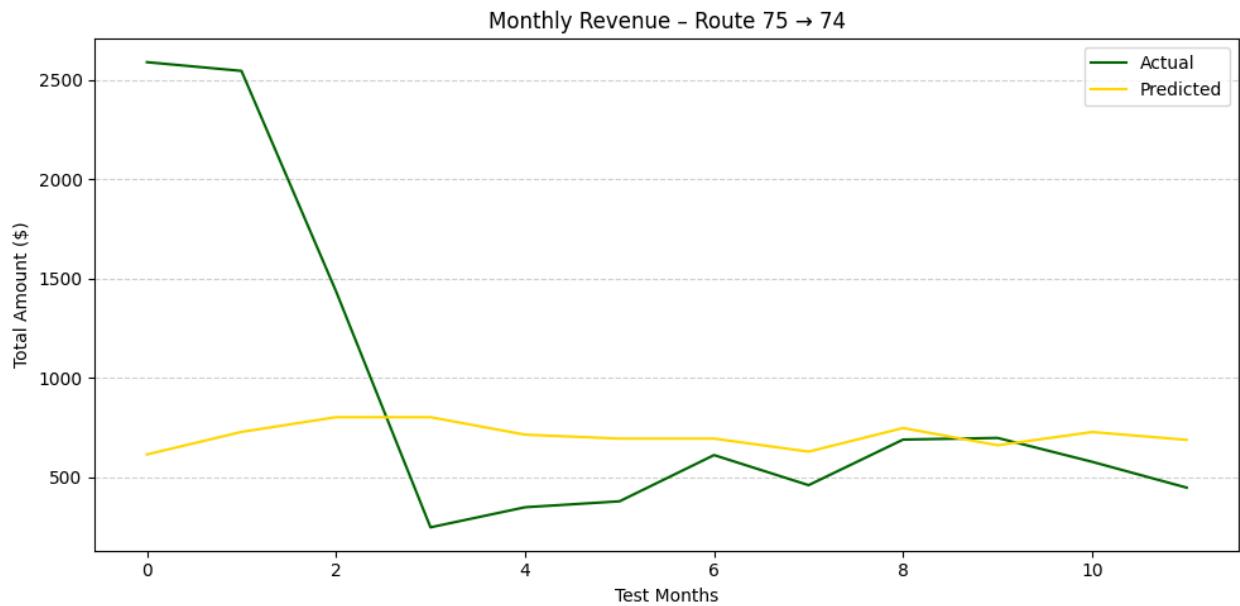
Using Route: PULocationID=75 → DOLocationID=74  
Total records for selected route: 4594

## Monthly revenue preview:

	year_month	monthly_revenue
0	2017-01-01	706.56
1	2017-02-01	716.12
2	2017-03-01	815.45
3	2017-04-01	655.81
4	2017-05-01	736.07

□ Route Revenue Forecasting Results

RMSE: 829.7019772888709  
R<sup>2</sup> Score: -0.09976580535236113



```
taxi_zones = pd.read_csv(  
    "C:/Users/parth/OneDrive/Desktop/Trips 2/taxi_zones.csv")  
  
df = df.merge(  
    taxi_zones[['LocationID', 'Borough']],  
    left_on='PULocationID',  
    right_on='LocationID',  
    how='left'  
)  
  
df = df.rename(columns={'Borough': 'pickup_borough'})  
df = df.drop(columns=['LocationID'])  
  
df = df.merge(  
    taxi_zones[['LocationID', 'Borough']],  
    left_on='DOLocationID',  
    right_on='LocationID',  
    how='left'  
)  
  
df = df.rename(columns={'Borough': 'dropoff_borough'})  
df = df.drop(columns=['LocationID'])  
  
df.columns.tolist()  
['VendorID',  
 'lpep_pickup_datetime',  
 'lpep_dropoff_datetime',
```

```
'store_and_fwd_flag',
'RatecodeID',
'PULocationID',
'DOLocationID',
'passenger_count',
'trip_distance',
'fare_amount',
'extra',
'mta_tax',
'tip_amount',
'tolls_amount',
'improvement_surcharge',
'total_amount',
'payment_type',
'trip_type',
'congestion_surcharge',
'week_start',
'pickup_borough',
'dropoff_borough',
'dropoff_borough',
'dropoff_borough',
'pickup_borough',
'dropoff_borough']

# Keep only first occurrence of each column name
df = df.loc[:, ~df.columns.duplicated()]

df.columns.tolist() # just to check, you should now see each name
only once

['VendorID',
'lpep_pickup_datetime',
'lpep_dropoff_datetime',
'store_and_fwd_flag',
'RatecodeID',
'PULocationID',
'DOLocationID',
'passenger_count',
'trip_distance',
'fare_amount',
'extra',
'mta_tax',
'tip_amount',
'tolls_amount',
'improvement_surcharge',
'total_amount',
'payment_type',
'trip_type',
'congestion_surcharge',
'week_start',
```

```

'pickup_borough',
'dropoff_borough']

df['route'] = df['pickup_borough'] + " → " + df['dropoff_borough']
df[['pickup_borough', 'dropoff_borough', 'route']].head()

C:\Users\parth\AppData\Local\Temp\ipykernel_15004\3927191911.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df['route'] = df['pickup_borough'] + " → " + df['dropoff_borough']

  pickup_borough dropoff_borough          route
0      Brooklyn      Brooklyn  Brooklyn → Brooklyn
1      Brooklyn      Brooklyn  Brooklyn → Brooklyn
2   Manhattan    Manhattan  Manhattan → Manhattan
3      Queens        Queens  Queens → Queens
4      Brooklyn      Brooklyn  Brooklyn → Brooklyn

df.loc[:, 'route'] = df['pickup_borough'] + " → " +
df['dropoff_borough']

# 1. Filter to years 2017–2020 (if not already done)

df['lpep_pickup_datetime'] =
pd.to_datetime(df['lpep_pickup_datetime'], errors='coerce')
df = df[(df['lpep_pickup_datetime'].dt.year >= 2017) &
        (df['lpep_pickup_datetime'].dt.year <= 2020)]

# 2. Check top borough→borough routes

print("Top 10 borough routes by trip count:")
print(df['route'].value_counts().head(10))

# Pick one route to forecast (you can change this string)
target_route = df['route'].value_counts().index[0]  # or e.g.
"Manhattan → Brooklyn"
print("\nUsing route:", target_route)

route_df = df[df['route'] == target_route].copy()

# =====
# 3. Build monthly revenue series for that route
# =====
route_df['year_month'] =
route_df['lpep_pickup_datetime'].dt.to_period('M').dt.to_timestamp()

```

```

monthly_rev = (route_df
               .groupby('year_month')['total_amount']
               .sum()
               .reset_index(name='monthly_revenue')
               .sort_values('year_month'))

print("\nMonthly revenue preview:")
print(monthly_rev.head())

# 4. Create lag features

monthly_rev['lag_1'] = monthly_rev['monthly_revenue'].shift(1)
monthly_rev['lag_2'] = monthly_rev['monthly_revenue'].shift(2)

# Drop first rows with NaN lags
monthly_ml = monthly_rev.dropna().reset_index(drop=True)

X = monthly_ml[['lag_1', 'lag_2']]
y = monthly_ml['monthly_revenue']

# 5. Train / test split
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, shuffle=False
)

model = RandomForestRegressor(
    n_estimators=150,
    random_state=42
)

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

rmse = mean_squared_error(y_test, y_pred) ** 0.5
r2 = r2_score(y_test, y_pred)

print("\nRoute Revenue Forecasting Results (borough based)")
print("Route:", target_route)
print("RMSE:", rmse)
print("R²:", r2)

# 6. Plot Actual vs Predicted monthly revenue

```

```
plt.figure(figsize=(10,5))
plt.plot(y_test.values, label="Actual", color="#006400")      # dark
green
plt.plot(y_pred, label="Predicted", color="#FFD700")        # dark
yellow
plt.title(f"Monthly Revenue - {target_route}")
plt.xlabel("Test Months")
plt.ylabel("Total Amount")
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

```
C:\Users\parth\AppData\Local\Temp\ipykernel_15004\2662573679.py:4:
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df['lpep_pickup_datetime'] =
pd.to_datetime(df['lpep_pickup_datetime'], errors='coerce')
```

```
Top 10 borough routes by trip count:
```

```
route
Manhattan → Manhattan    119029
Queens → Queens           97614
Brooklyn → Brooklyn       97315
Bronx → Bronx              19023
Brooklyn → Manhattan      17506
Manhattan → Bronx          11054
Queens → Manhattan         7588
Brooklyn → Queens           6708
Bronx → Manhattan           6047
Queens → Brooklyn            5205
Name: count, dtype: int64
```

```
Using route: Manhattan → Manhattan
```

```
Monthly revenue preview:
```

```
year_month  monthly_revenue
0 2017-01-01      29748.80
1 2017-02-01      29763.81
2 2017-03-01      31136.45
3 2017-04-01      29900.01
4 2017-05-01      30307.12
```

```
□ Route Revenue Forecasting Results (borough based)
```

```
Route: Manhattan → Manhattan
```

RMSE: 30430.07166554574  
R<sup>2</sup>: -0.017662602963947505

